

Exceptional service in the national interest



Sandia
National
Laboratories



docker Containers & Interactive Environments

Using Trilinos from

Tobias Wiesner

May 12, 2019



Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy a National Nuclear Security Administration under contract DE-NA0003525. SAND NO. 2016-10805 C

Typical workflow from the first idea to the final performance tests:

1) Idea & Algorithm design

- Rapid development (Matlab, Python, C++)
- Quick experiments on small data sets

2) Implementation

- Final parallelized implementation in C++ (using Trilinos)
- Algorithmic optimizations & Performance optimizations

3) Performance studies on HPC clusters

- Final performance studies on HPC clusters with large datasets
- Produce results for scientific paper

What are the challenges?

1) Idea & Algorithm design

- Focus on algorithms and mathematics
- Need for rapid development tools and programming languages

2) Implementation

- Error-prone reimplementing and parallelization of algorithms using Trilinos
- Expert knowledge for best results necessary (Trilinos, etc.)

3) Performance studies on HPC clusters

- Make your algorithms work on all HPC platforms

New technologies and tools

What is docker?

Docker is a software containerization platform.

- New concepts for *organizing, developing, testing* and *shipping* software
- Maintain one general software/development environment independent of underlying host OS or hardware
- **Focus on the applications!**

More information: <https://docs.docker.com/install/>

What is jupyter?

Jupyter is an interactive computing environment.

- Support for many language kernels including Python, Julia, R, C++
- Ideal tool for interactive tutorials
- Jupyter notebooks for documenting and organizing scientific experiments

- Document your ideas, algorithms, experiments and results in notebooks
- Use functionality from existing C++ libraries in your notebooks
- Extend the functionality by writing new classes and functions live in your notebook
- Gain better insight into legacy code and complex algorithms by adding advanced debugging statements
- Optimize your algorithms in an interactive way

Example

`AutoWhiteBalancingLiveDemo.html`

What jupyter is and what it is NOT

Jupyter

- helps you minimizing the effort for porting the Jupyter C++ implementation to the final implementation.
- helps you developing algorithms and gaining in-depth insight.
- helps you organizing your ideas, experiments and results.

Jupyter does not

- make porting the experimental C++ code to final C++ code superfluous.
- help you writing clean code and unit tests.
- help you organizing your notebooks. You still need a systematic way of storing your notebooks in order to quickly find results from a specific experiment. Use a version controlling system!

New technologies in context of Trilinos

How can docker and jupyter help?

1) Idea & Algorithm design

- Use C++ with Cling and Jupyter as rapid prototyping environment?

2) Implementation

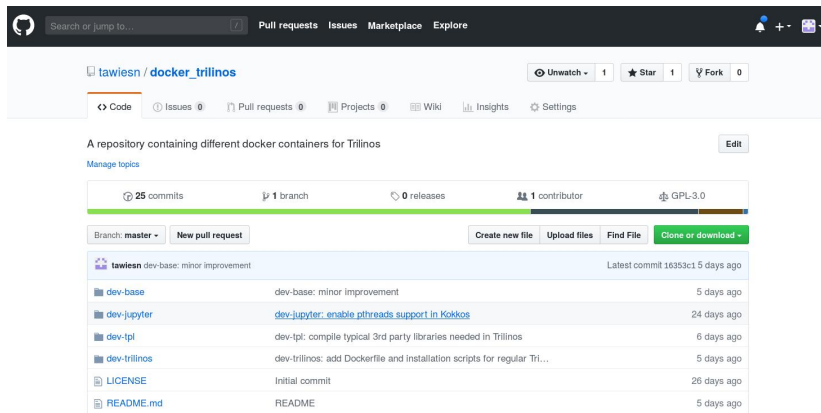
- Improve your Trilinos knowledge with Jupyter Trilinos tutorials
- Manage your development environment with docker

3) Performance studies on HPC clusters

- Manage your development environment with docker

Check out the github repository

https://github.com/tawiesn/docker_trilinos



The screenshot shows the GitHub repository page for `tawiesn/docker_trilinos`. The repository is described as "A repository containing different docker containers for Trilinos". It has 25 commits, 1 branch, 0 releases, 1 contributor, and is licensed under GPL-3.0. The repository is currently on the `master` branch. The commit history shows several recent commits, including `dev-base: minor improvement`, `dev-jupyter: enable pthreads support in Kokkos`, `dev-tpi: compile typical 3rd party libraries needed in Trilinos`, and `dev-trilinos: add Dockerfile and installation scripts for regular Tri...`.

Search or jump to...

Pull requests Issues Marketplace Explore

tawiesn / docker_trilinos

Unwatch 1 Star 1 Fork 0

Code Issues 0 Pull requests 0 Projects 0 Wiki Insights Settings

A repository containing different docker containers for Trilinos

Manage topics

25 commits 1 branch 0 releases 1 contributor GPL-3.0

Branch: master New pull request

Create new file Upload files Find File Clone or download

tawiesn dev-base: minor improvement		Latest commit 16353c1 5 days ago
dev-base	dev-base: minor improvement	5 days ago
dev-jupyter	dev-jupyter: enable pthreads support in Kokkos	24 days ago
dev-tpi	dev-tpi: compile typical 3rd party libraries needed in Trilinos	6 days ago
dev-trilinos	dev-trilinos: add Dockerfile and installation scripts for regular Tri...	5 days ago
LICENSE	Initial commit	26 days ago
README.md	README	5 days ago

Docker for Trilinos

The dev-trilinos docker container

- Classical docker container based on CentOS 7 (Redhat compatible)
- Contains a clone of the public Trilinos repository
- Builds some core Trilinos packages using gcc with OpenMPI support
- **Meant as an example for your own Trilinos container** that can be used to link your application against
- Use it as **standardized development platform** that can also be used on **HPC clusters**, e.g. via the Shifter project (<https://github.com/NERSC/shifter>)

More information here: https://github.com/tawiesn/docker_trilinos/blob/master/README.md

How to create your own Trilinos container for your application



- Clone the `https://github.com/tawiesn/docker_trilinos` repository
- Follow the instructions to build all dependencies. For details see `https://github.com/tawiesn/docker_trilinos/blob/master/installation_dev_trilinos.md`
- Create a copy of the dev-trilinos folder and rename it, e.g. my-dev-trilinos
- Edit the my-dev-trilinos/do-configure file and enable all Trilinos features that you need
- If you need a specific Trilinos version for your application you might adapt the git clone statement in the my-dev-trilinos/Dockerfile

How to create your own application docker container

- Build your my-dev-trilinos docker container
- Create a new folder, e.g. my-application
- Create a new Dockerfile in you my-application folder
- Edit the my-application/Dockerfile and make sure you derive your docker container from the my-dev-trilinos folder.
- Copy the sources and build scripts into the docker container. Use the dev-trilinos/Dockerfile for some inspiration.

Benefit

Your Dockerfiles contain all necessary commands to compile and install your Trilinos libraries and application.

- Put your Dockerfiles under version control
- Set up nightly builds for the Docker container to track changes and detect compilation issues as soon as they pop up

Jupyter and Trilinos

The dev-jupyter docker container

- Docker container based on CentOS 7 (Redhat compatible) with clang installation and cling extensions
- Contains a clone of the public Trilinos repository
- **Meant for tutorials and algorithmic experiments**
- Find a small set of basic Trilinos tutorial notebooks here:
<https://github.com/tawiesn/trilinos-notebooks>
- Limitations:
 - No OpenMPI support at the moment (can be added)
 - Some instabilities when running parallel Kokkos kernels (needs to be investigated)

More information here: https://github.com/tawiesn/docker_trilinos/blob/master/README.md

Repository

<https://github.com/tawiesn/trilinos-notebooks>

1) Epetra tutorial

Apply power method to tridiagonal matrix

2) Belos tutorial

Solve linear equation with GMRES w/wo ILU_t

3) Tpetra tutorial

Create Tpetra vector and access data

4) Kokkos tutorial

Some basic examples for parallel kernels and Kokkos views

For more information check out the repositories:

Docker container



[https://github.com/
tawiesn/docker_trilinos](https://github.com/tawiesn/docker_trilinos)

Jupyter notebooks



[https://github.com/
tawiesn/trilinos-notebooks](https://github.com/tawiesn/trilinos-notebooks)